



PART 8 — Backend Python 🚀

The layer where Python becomes production backend engineering

This part is **very important for your FastAPI work** because these concepts connect your Python fundamentals to the actual architecture of a production API.






FastAPI Clean Architecture - Layer Separation



FastAPI + SQLAlchemy + PostgreSQL: How It Actually Works


Sessions. Transactions. Migrations. Production Architecture.






MASTER MENTAL MODEL

Think of a backend application as a pipeline:





This is the **big picture** you should have in your head.

1. Virtual Environment

A virtual environment gives each project its own isolated Python environment.

```
python -m venv .venv
```

Then:

```
source .venv/bin/activate
```

Windows:

```
.venv\Scripts\activate
```

Conceptually:

```
System Python
├── Project A
│   └── .venv
│       ├── FastAPI 0.x
│       └── SQLAlchemy 2.x
└── Project B
    └── .venv
        ├── Django
        └── SQLAlchemy another version
```

Without isolation:

```
Everything
↓
Global Python
↓
Dependency conflicts 💣
```

With isolation:

```
Project
  ↓
.venv
  ↓
Independent dependencies
```

Python's `venv` creates lightweight environments with their own installed packages and interpreter-related configuration. Environments are generally disposable and should not be committed to Git. ([Python documentation](#))

Senior interview answer

A virtual environment isolates a project's Python interpreter and dependencies so projects don't interfere with each other.

2. `pip`

`pip` is the package installer.

```
pip install fastapi
```

Better practice:

```
python -m pip install fastapi
```

Why?

Because this explicitly associates `pip` with the Python interpreter you're invoking.

Typical:

```
python -m pip install fastapi
python -m pip install sqlalchemy
python -m pip uninstall requests
python -m pip list
```

What happens during install?

Conceptually:

```
pip install fastapi
  ↓
Find package
  ↓
Read dependencies
  ↓
Resolve compatible versions
  ↓
Download/build distributions
  ↓
Install
```

pip performs dependency resolution and may backtrack when package constraints conflict. ([Pip](#))

3. requirements.txt

A `requirements.txt` is essentially:

A list of packages/requirements that pip should install.

Example:

```
fastapi
uvicorn
sqlalchemy
asyncpg
pydantic
pytest
```

Then:


```
pip install -r requirements.txt
```

You can specify versions:

```
fastapi==0.116.1
sqlalchemy>=2.0,<3.0
```

Visual:

```
requirements.txt
|
```

pip's requirements-file format is specifically designed to provide items for `pip install`; despite the common filename, the file need not literally be named `requirements.txt`. (Pip)

4. `pyproject.toml` ★★★★★

This is much more important for modern Python projects.

Think:

```
requirements.txt
  ↓
Mostly dependency installation list

pyproject.toml
  ↓
Project configuration / metadata
  +
Dependencies
  +
Build configuration
  +
Tool configuration
```

Example:

```
[project]
name = "my-api"
version = "1.0.0"
requires-python = ">=3.12"

dependencies = [
    "fastapi",
    "sqlalchemy",
    "asyncpg",
    "pydantic",
]

[project.optional-dependencies]
test = [
```

```
"pytest",  
"httpx",  
]
```

And:

```
[tool.pytest.ini_options]  
testpaths = ["tests"]
```

Modern `pyproject.toml` can contain:

```
[build-system]  
[project]  
[tool.*]
```

The `[project]` table contains standardized project metadata and dependencies, while `[tool]` is for tool-specific configuration. ([Python Packaging](#))

Senior interview distinction

```
requirements.txt  
  ↓  
pip-oriented dependency list  
  
pyproject.toml  
  ↓  
Modern project definition  
+ metadata  
+ dependencies  
+ build system  
+ tool configuration
```

For a new project, **know `pyproject.toml` very well.**

5. Type Hints

You already learned the basics.

Now understand why backend developers care.

Without:

```
def get_user(id):  
    ...
```

With:

```
def get_user(id: int) -> User:  
    ...
```

Now the code communicates:

```
id  
↓  
int  
  
return  
↓  
User
```

Type hints help:

- IDEs
- static type checkers
- linters
- refactoring
- API frameworks
- documentation
- maintainability

Python itself does **not** enforce annotations at runtime. They are primarily information used by tools unless a framework/library explicitly uses them. ([Python documentation](#))

★ Why FastAPI LOVES Type Hints

Consider:

```
@app.get("/users/{user_id}")  
async def get_user(user_id: int):  
    ...
```

FastAPI can use that declaration to understand:

```
user_id
  ↓
integer
  ↓
validate
  ↓
OpenAPI documentation
```

FastAPI uses Python type declarations for parameter requirements, conversion, validation and API documentation. ([FastAPI](#))

That's one of the biggest reasons FastAPI feels so natural to Python developers.

6. Pydantic

Pydantic is about:

Data validation + parsing + serialization using Python type declarations.

Example:

```
from pydantic import BaseModel

class UserCreate(BaseModel):
    name: str
    age: int
    email: str
```

Incoming:

```
{
  "name": "John",
  "age": 30,
  "email": "john@example.com"
}
```

Pydantic:

```
Raw Input
  ↓
Pydantic Model
  ↓
Validation
```

↓
Validated Python object

If invalid:

Raw Input
↓
Pydantic
↓
❌ ValidationError

Pydantic models are classes derived from `BaseModel` whose fields are declared as annotated attributes. Current Pydantic supports validation, serialization, JSON Schema generation, strict/lax modes and custom validators. ([Pydantic](#))

Pydantic + FastAPI

This is a critical relationship:

```
class UserCreate(BaseModel):  
    name: str  
    age: int  
  
@app.post("/users")  
async def create_user(user: UserCreate):  
    return user
```

Request:

HTTP JSON
↓
FastAPI
↓
Pydantic
↓
UserCreate
↓
Endpoint

Response:

```
Python object
  ↓
Pydantic / FastAPI
  ↓
JSON
```

FastAPI is built around Python type hints and Pydantic for data handling. ([FastAPI](#))

7. FastAPI 🌟🌟🌟🌟🌟

Now connect everything.

FastAPI is your **HTTP application framework**.

Example:

```
from fastapi import FastAPI

app = FastAPI()

@app.get("/users/{user_id}")
async def get_user(user_id: int):
    return {"id": user_id}
```

Architecture:

```
Client
  |
  | HTTP
  ↓
FastAPI
  |
  ├── Routing
  ├── Validation
  ├── Dependency Injection
  ├── Middleware
  ├── Serialization
  └── OpenAPI
      |
      ↓
      Endpoint
```

The important thing is:

FastAPI isn't your business logic.

Good architecture:

```
Router
  ↓
Service
  ↓
Repository
  ↓
Database
```

Instead of:

```
Router
  ↓
500 lines of business logic 🤖
  ↓
Database
```

8. Dependency Injection ★★★★★

This is one of the **most important FastAPI concepts**.

Suppose your endpoint needs a database session.

Instead of:

```
async def endpoint():
    session = create_session()
    ...
    await session.close()
```


you can say:

```
async def endpoint(
    session: Annotated[AsyncSession, Depends(get_session)]
):
    ...
```

You're telling FastAPI:

"I need an `AsyncSession`. You figure out how to provide it."

Visual:



FastAPI resolves dependencies and injects their results into the endpoint. Dependencies can themselves have dependencies, forming a dependency graph. ([FastAPI](#))

Dependency Injection Example

```
async def get_db():
    async with SessionLocal() as session:
        yield session

@app.get("/users")
async def get_users(
    db: Annotated[AsyncSession, Depends(get_db)]
):
    result = await db.execute(...)
    return result.scalars().all()
```

Mental model:


```
Request
↓
FastAPI
↓
get_db()
↓
AsyncSession
↓
Endpoint
↓
Response
```

↓
cleanup

This is why **yield** dependencies are particularly powerful for database sessions and resource cleanup. FastAPI uses context-manager machinery internally for this pattern. ([FastAPI](#))

9. Middleware ★★★★★

Middleware surrounds the request/response cycle.



Example:

```
@app.middleware("http")
async def log_requests(request, call_next):
    print(request.url)

    response = await call_next(request)

    print(response.status_code)

    return response
```

Typical middleware responsibilities:

- request logging
- request IDs
- CORS

- authentication-related processing
- timing
- metrics
- headers
- tracing

FastAPI middleware receives the request before the route handler and can also process the response before it is returned. ([FastAPI](#))

Dependency vs Middleware

```
Middleware
  ↓
Usually application-wide request/response concern

Dependency
  ↓
Specific resource/business requirement
```

Example:

```
Middleware → request ID
Dependency → current user
Dependency → database session
Middleware → timing
```

10. Background Tasks

Sometimes the client doesn't need to wait for work.

Example:

```
POST /send-email
  ↓
Save request
  ↓
Return 200
  ↓
Send email
```


FastAPI:

```
from fastapi import BackgroundTasks

@app.post("/notify")
async def notify(background_tasks: BackgroundTasks):
    background_tasks.add_task(send_email)

    return {"status": "accepted"}
```

Visual:



FastAPI's **BackgroundTasks** are intended for work that can happen after the response is returned, such as small notification or processing jobs. For heavy/distributed workloads, FastAPI's documentation recommends considering a job system such as Celery instead. ([FastAPI](#))

Important distinction

Don't use:

BackgroundTasks

for:

10-hour ML computation
Huge video processing
Large distributed workloads

Use a proper worker/queue architecture.

11. SQLAlchemy ★★★★★

SQLAlchemy is your database toolkit/ORM.

Without ORM:

```
SELECT *  
FROM users  
WHERE id = 10;
```

With SQLAlchemy:

```
stmt = select(User).where(User.id == 10)
```

Think:

```
Python  
↓  
SQLAlchemy  
↓  
SQL  
↓  
Database
```

SQLAlchemy has two major conceptual layers:

```
SQLAlchemy  
├── Core  
│   └── SQL / database toolkit  
└── ORM  
    └── Python objects ↔ database tables
```

Example ORM:

```
class User(Base):  
    __tablename__ = "users"  
  
    id: Mapped[int] = mapped_column(primary_key=True)  
    name: Mapped[str]
```

Database:

users

id	name
1	John
2	Alex

Python:

```
User(id=1, name="John")
```

12. Async SQLAlchemy

This is directly relevant to your FastAPI work.

Typical setup:

```
from sqlalchemy.ext.asyncio import (
    AsyncSession,
    async_sessionmaker,
    create_async_engine,
)

engine = create_async_engine(DATABASE_URL)

SessionLocal = async_sessionmaker(
    engine,
    expire_on_commit=False,
)
```

Then:

```
async with SessionLocal() as session:
    result = await session.execute(
        select(User)
    )

    users = result.scalars().all()
```

Architecture:

```
FastAPI
  |
  ↓
async endpoint
  |
  ↓
AsyncSession
  |
  ↓
Async DB Driver
  |
  ↓
Database
```

SQLAlchemy's asyncio extension provides `AsyncEngine`, `AsyncConnection`, and `AsyncSession`; current SQLAlchemy documentation specifically warns that a single `AsyncSession` is not safe to share across concurrent asyncio tasks. (docs.sqlalchemy.org)



Critical Async SQLAlchemy Rule

Don't do:

```
tasks = [
    asyncio.create_task(query_1(session)),
    asyncio.create_task(query_2(session)),
]
```

with the **same** `AsyncSession`.

Instead:

```
Task A → Session A
Task B → Session B
Task C → Session C
```

SQLAlchemy explicitly documents `AsyncSession` as mutable/stateful and unsafe for concurrent use by multiple asyncio tasks. (docs.sqlalchemy.org)

This is a **very good senior interview question**.

13. pytest ★★★★★

pytest is your testing framework.

Simple:

```
def add(a, b):  
    return a + b
```

Test:

```
def test_add():  
    assert add(2, 3) == 5
```

Run:

```
pytest
```

Mental model:

```
Application  
  ↓  
pytest  
  ↓  
Tests  
  ↓  
Assertions  
  ↓  
PASS / FAIL
```

FastAPI Testing

FastAPI provides **TestClient**:

```
from fastapi.testclient import TestClient

client = TestClient(app)

def test_users():
    response = client.get("/users")

    assert response.status_code == 200
```

FastAPI's testing support uses HTTPX/Starlette's `TestClient` and integrates naturally with pytest. ([FastAPI](#))

Fixtures ★★★★★

One of pytest's most important concepts:

```
@pytest.fixture
def user():
    return User(
        id=1,
        name="John"
    )
```

Then:

```
def test_user(user):
    assert user.name == "John"
```

Mental model:

```
Fixture
↓
Prepare dependency/test data
↓
Test
↓
Cleanup
```

Fixtures become extremely important when testing:

- database sessions

- authentication
- test users
- mock services
- API clients
- temporary files

14. Logging

Never build production backend observability around:

```
print("something happened")
```

Use:

```
import logging

logger = logging.getLogger(__name__)

logger.info("User created")
logger.warning("Slow request")
logger.error("Database failure")
```

Python logging has standard levels:

```
DEBUG
↓
INFO
↓
WARNING
↓
ERROR
↓
CRITICAL
```

Python's logging system is built around loggers, handlers, filters and formatters. The recommended module-level pattern is:

```
logger = logging.getLogger(__name__)
```

rather than using the root logger directly. ([Python documentation](#))

Production Logging

You want something like:

```
2026-09-02 21:10:31
INFO
request_id=abc123
user_id=42
GET /users/42
200
85ms
```

Not:

```
hello
user called endpoint
something happened
```

Good backend logs should help answer:

```
What happened?
When?
Where?
For which request?
For which user?
How long?
Did it fail?
Why?
```

15. Configuration / `.env` ★★★★★

Never hard-code secrets:

```
DATABASE_URL = "postgresql://admin:password123@..."
```

Instead:

```
Environment
↓
.env
```

```
↓  
Configuration  
↓  
Application
```

Example `.env`:

```
DATABASE_URL=postgresql://localhost/mydb  
DEBUG=false  
API_KEY=secret
```

Then your configuration layer reads them.


For example with Pydantic Settings:

```
from pydantic_settings import BaseSettings  
  
class Settings(BaseSettings):  
    database_url: str  
    debug: bool = False  
    api_key: str
```

Then:

```
settings = Settings()
```

Conceptually:



Important


`.env` is convenient for local development.


Do **not** commit secrets to Git.

Production environments commonly inject secrets through the deployment/platform's secret-management mechanism.

🔥 Complete Production Request Flow

Now combine **all 15 concepts**:





pytest
tests the entire application

Logging
observes the entire system

Configuration
.env / environment

Backend Project Structure

A structure you should understand well:

```
project/
├── app/
│   ├── main.py
│   └── api/
│       └── routes/
│           ├── users.py
│           └── auth.py
```

```
├── schemas/
│   ├── user.py
│   └── auth.py
├── models/
│   └── user.py
├── services/
│   └── user_service.py
├── repositories/
│   └── user_repository.py
├── db/
│   ├── session.py
│   └── base.py
├── core/
│   ├── config.py
│   └── logging.py
├── dependencies/
│   └── common.py
├── tests/
│   ├── unit/
│   └── integration/
├── .env
├── .gitignore
├── pyproject.toml
└── README.md
```

Think of it as:

```
Router
  ↓
Schema
  ↓
Service
  ↓
Repository
  ↓
SQLAlchemy
  ↓
Database
```


★ The Most Important Distinctions

requirements.txt vs pyproject.toml

```
requirements.txt
→ pip installation requirements

pyproject.toml
→ project metadata + dependencies
→ build configuration
→ tool configuration
```

Type Hint vs Pydantic

```
Type Hint
→ describes expected type

Pydantic
→ validates/parses data against a schema
```

Example:

```
def create_user(age: int):
    ...
```

vs:

```
class User(BaseModel):
    age: int
```

Middleware vs Dependency

```
Middleware
→ surrounds request/response pipeline
```

Dependency
→ provides something an endpoint depends on

SQLAlchemy vs Pydantic

This is **extremely important**.

Pydantic
→ API/data validation
→ request/response schemas

SQLAlchemy
→ database interaction
→ ORM/database models

Don't automatically use one model for everything.

Typical architecture:


```
JSON
↓
Pydantic Request Model
↓
Service
↓
SQLAlchemy ORM Model
↓
Database

Database
↓
SQLAlchemy ORM Model
↓
Pydantic Response Model
↓
JSON
```

Senior-Level Architecture

A mature FastAPI application often looks conceptually like:

HTTP
|



This separation is what keeps a backend maintainable as it grows.

PART 8 CHEAT SHEET

Concept	Mental model
Virtual Environment	Isolated Python environment
<code>pip</code>	Package installer/dependency resolver
<code>requirements.txt</code>	pip-oriented dependency list
<code>pyproject.toml</code>	Modern project/build/tool configuration
Type Hints	Describe expected types
Pydantic	Validate/parse/serialize structured data
FastAPI	Web/API framework
Dependency Injection	Framework provides required dependencies

Concept	Mental model
Middleware	Code surrounding request/response
Background Tasks	Small work after response
SQLAlchemy	Database toolkit + ORM
Async SQLAlchemy	Async database access
pytest	Testing framework
Logging	Runtime observability
<code>.env</code>	Local/environment configuration source

🔥 10 Questions You MUST Master

If you're preparing for **senior Python + FastAPI interviews**, make sure you can explain these without memorizing definitions:

1.

Why use a virtual environment?

2.

`requirements.txt` vs `pyproject.toml`?

3.

Are Python type hints enforced at runtime?

4.

What does Pydantic actually do?

5.

How does FastAPI use type hints and Pydantic?

6.

How does FastAPI Dependency Injection work internally/conceptually?

7.

Middleware vs dependency — when do you use each?

8.

Why shouldn't you use the same `AsyncSession` across concurrent tasks?

9.


When should FastAPI **BackgroundTasks** be replaced by a real job queue?


10.

Why separate Pydantic schemas from SQLAlchemy ORM models?



Your Backend Master Map





The core progression is:

**Python → Type Hints → Pydantic → FastAPI → Dependency Injection → SQLAlchemy → Async
SQLAlchemy → Testing → Logging → Production Configuration**

Once this becomes second nature, you're no longer just learning Python—you are understanding the **architecture of a production Python backend**.